

# fct0vldxf

February 17, 2026

Task 03: Customer Prediction Dataset (Bank Customers).

Objective: Identify customers who are likely to leave the bank.

Dataset: Churn Modleing Dataset. About Dataset: Content: This data set contains details of a bank's customers and the target variable is a binary variable reflecting the fact whether the customer left the bank (closed his account) or he continues to be a customer.

Load the Dataset:

```
[1]: import pandas as pd
```

```
[2]: Dataset = pd.read_csv("Churn_Modelling.csv")
```

Data Exploration:

```
[3]: Dataset.head()
```

```
[3]:   RowNumber CustomerId Surname CreditScore Geography Gender Age \
0           1  15634602 Hargrave       619    France Female  42
1           2  15647311     Hill        608    Spain Female  41
2           3  15619304    Onio        502    France Female  42
3           4  15701354    Boni        699    France Female  39
4           5  15737888  Mitchell       850    Spain Female  43

      Tenure    Balance NumOfProducts HasCrCard IsActiveMember \
0         2      0.00            1          1             1
1         1    83807.86            1          0             1
2         8   159660.80            3          1             0
3         1      0.00            2          0             0
4         2   125510.82            1          1             1

  EstimatedSalary  Exited
0      101348.88      1
1      112542.58      0
2      113931.57      1
3      93826.63      0
4      79084.10      0
```

As we see in the data, fields like RowNumber, CustomerId and Surname do not matter at all during

out analysis. They are just additional infomration about the bank customers. So, we go ahead and drop these three columns.

[4]: `Dataset.drop(["RowNumber", "CustomerId", "Surname"], axis = 1, inplace = True)`

[5]: `Dataset.head()`

[5]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	\
0	619	France	Female	42	2	0.00		1
1	608	Spain	Female	41	1	83807.86		1
2	502	France	Female	42	8	159660.80		3
3	699	France	Female	39	1	0.00		2
4	850	Spain	Female	43	2	125510.82		1

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	1	101348.88	1
1	0	1	112542.58	0
2	1	0	113931.57	1
3	0	0	93826.63	0
4	1	1	79084.10	0

[6]: `Dataset.tail()`

[6]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	\
9995	771	France	Male	39	5	0.00		2
9996	516	France	Male	35	10	57369.61		1
9997	709	France	Female	36	7	0.00		1
9998	772	Germany	Male	42	3	75075.31		2
9999	792	France	Female	28	4	130142.79		1

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
9995	1	0	96270.64	0
9996	1	1	101699.77	0
9997	0	1	42085.58	1
9998	1	0	92888.52	1
9999	1	0	38190.78	0

[7]: `Dataset.shape`

[7]: `(10000, 11)`

[8]: `Dataset.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
---  -- 
 0   RowNumber       10000 non-null   int64 
 1   CreditScore     10000 non-null   float64
 2   Geography       10000 non-null   object 
 3   Gender          10000 non-null   object 
 4   Age              10000 non-null   int64 
 5   Tenure           10000 non-null   int64 
 6   Balance          10000 non-null   float64
 7   NumOfProducts    10000 non-null   int64 
 8   HasCrCard        10000 non-null   int64 
 9   IsActiveMember   10000 non-null   int64 
 10  EstimatedSalary 10000 non-null   float64
 11  Exited           10000 non-null   int64 
```

```
0    CreditScore      10000 non-null   int64
1    Geography        10000 non-null   object
2    Gender           10000 non-null   object
3    Age              10000 non-null   int64
4    Tenure           10000 non-null   int64
5    Balance          10000 non-null   float64
6    NumOfProducts    10000 non-null   int64
7    HasCrCard        10000 non-null   int64
8    IsActiveMember   10000 non-null   int64
9    EstimatedSalary  10000 non-null   float64
10   Exited           10000 non-null   int64
dtypes: float64(2), int64(7), object(2)
memory usage: 859.5+ KB
```

```
[9]: Dataset.describe()
```

```
[9]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	\
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	
mean	650.528800	38.921800	5.012800	76485.889288	1.530200	
std	96.653299	10.487806	2.892174	62397.405202	0.581654	
min	350.000000	18.000000	0.000000	0.000000	1.000000	
25%	584.000000	32.000000	3.000000	0.000000	1.000000	
50%	652.000000	37.000000	5.000000	97198.540000	1.000000	
75%	718.000000	44.000000	7.000000	127644.240000	2.000000	
max	850.000000	92.000000	10.000000	250898.090000	4.000000	

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	0.70550	0.515100	100090.239881	0.203700
std	0.45584	0.499797	57510.492818	0.402769
min	0.00000	0.000000	11.580000	0.000000
25%	0.00000	0.000000	51002.110000	0.000000
50%	1.00000	1.000000	100193.915000	0.000000
75%	1.00000	1.000000	149388.247500	0.000000
max	1.00000	1.000000	199992.480000	1.000000

```
[10]: Dataset.columns
```

```
[10]: Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',
       'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary',
       'Exited'],
       dtype='object')
```

```
[11]: Dataset.dtypes
```

```
[11]: CreditScore      int64
Geography         object
```

```
Gender          object
Age            int64
Tenure         int64
Balance        float64
NumOfProducts   int64
HasCrCard      int64
IsActiveMember int64
EstimatedSalary float64
Exited         int64
dtype: object
```

Data Visualization:

```
[12]: import numpy as np

import pandas as pd

import matplotlib as pyplot
from matplotlib import pyplot as plt

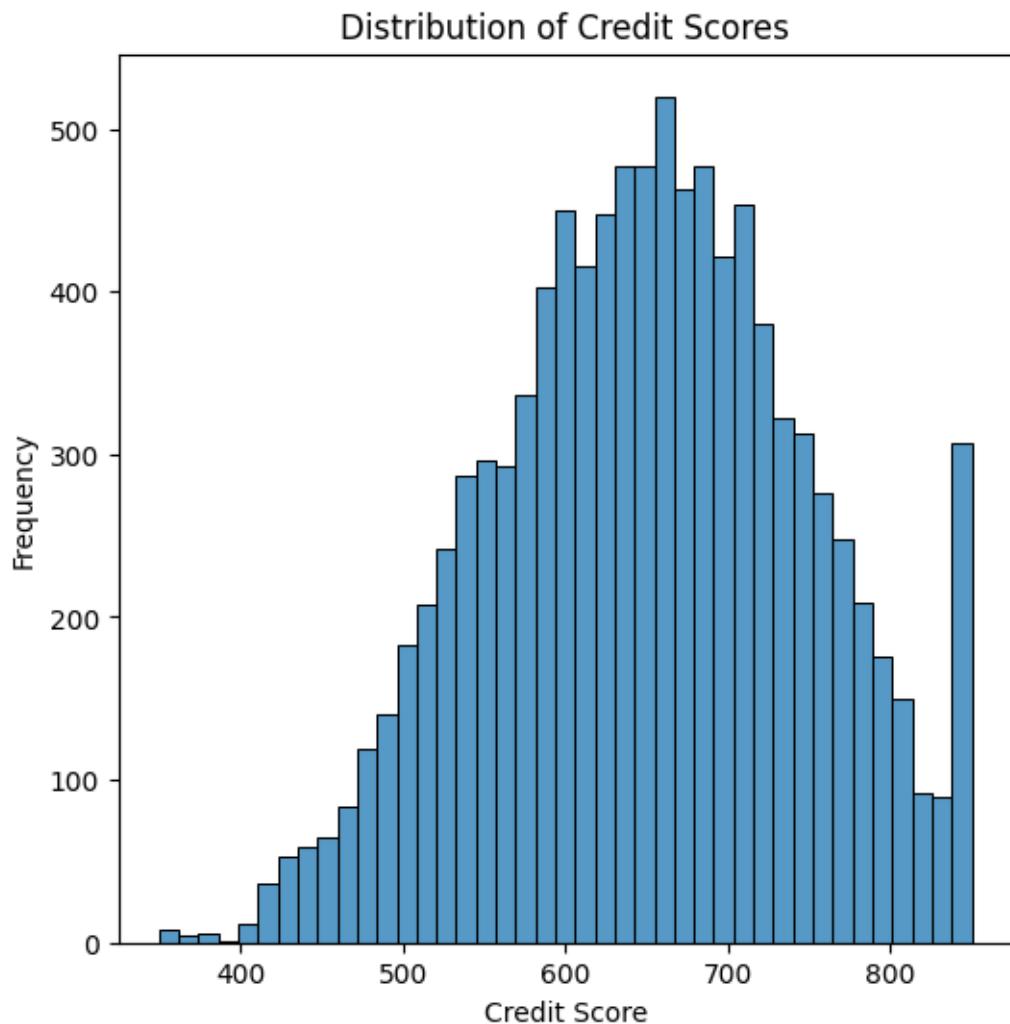
import seaborn as sns
```

```
[13]: plt.figure(figsize=(6,6))

sns.histplot(Dataset["CreditScore"])

plt.xlabel("Credit Score")
plt.ylabel("Frequency")
plt.title("Distribution of Credit Scores")

plt.show()
```

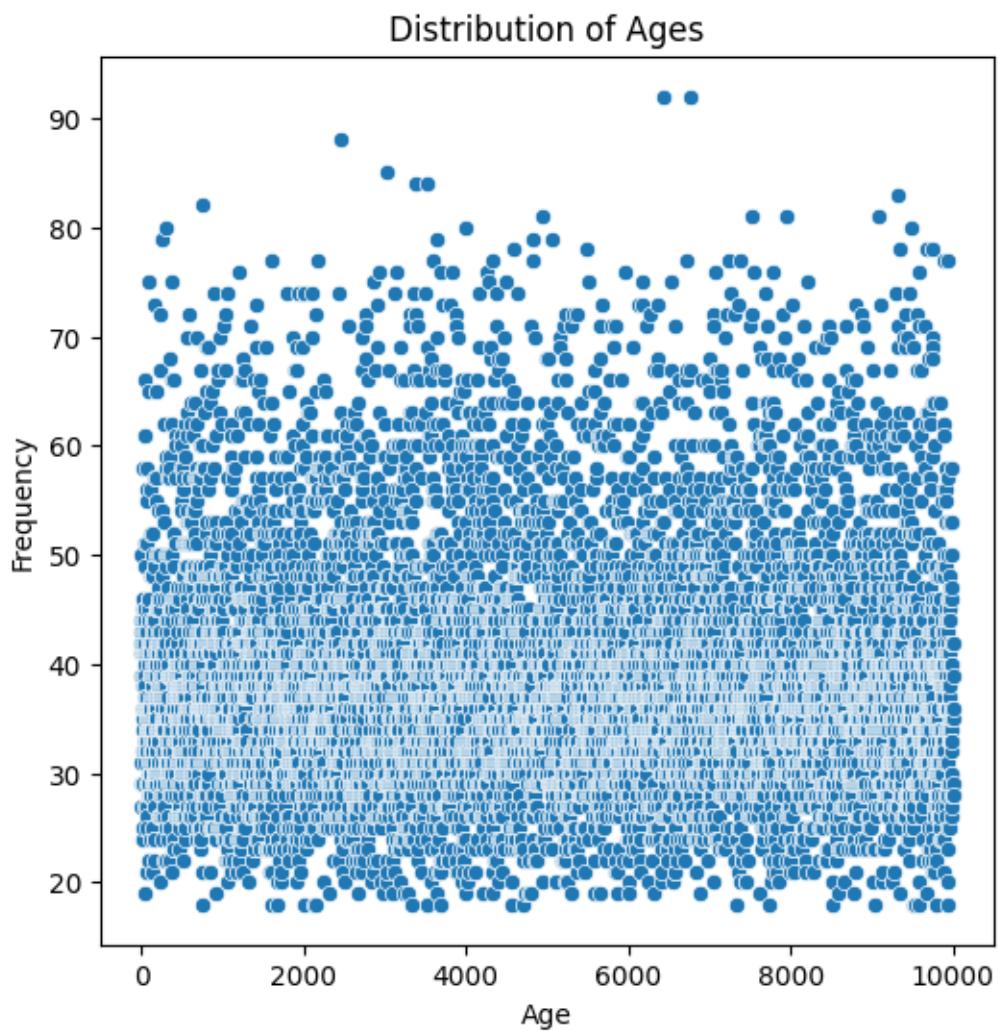


```
[14]: plt.figure(figsize=(6,6))

sns.scatterplot(Dataset["Age"])

plt.xlabel("Age")
plt.ylabel("Frequency")
plt.title("Distribution of Ages")

plt.show()
```

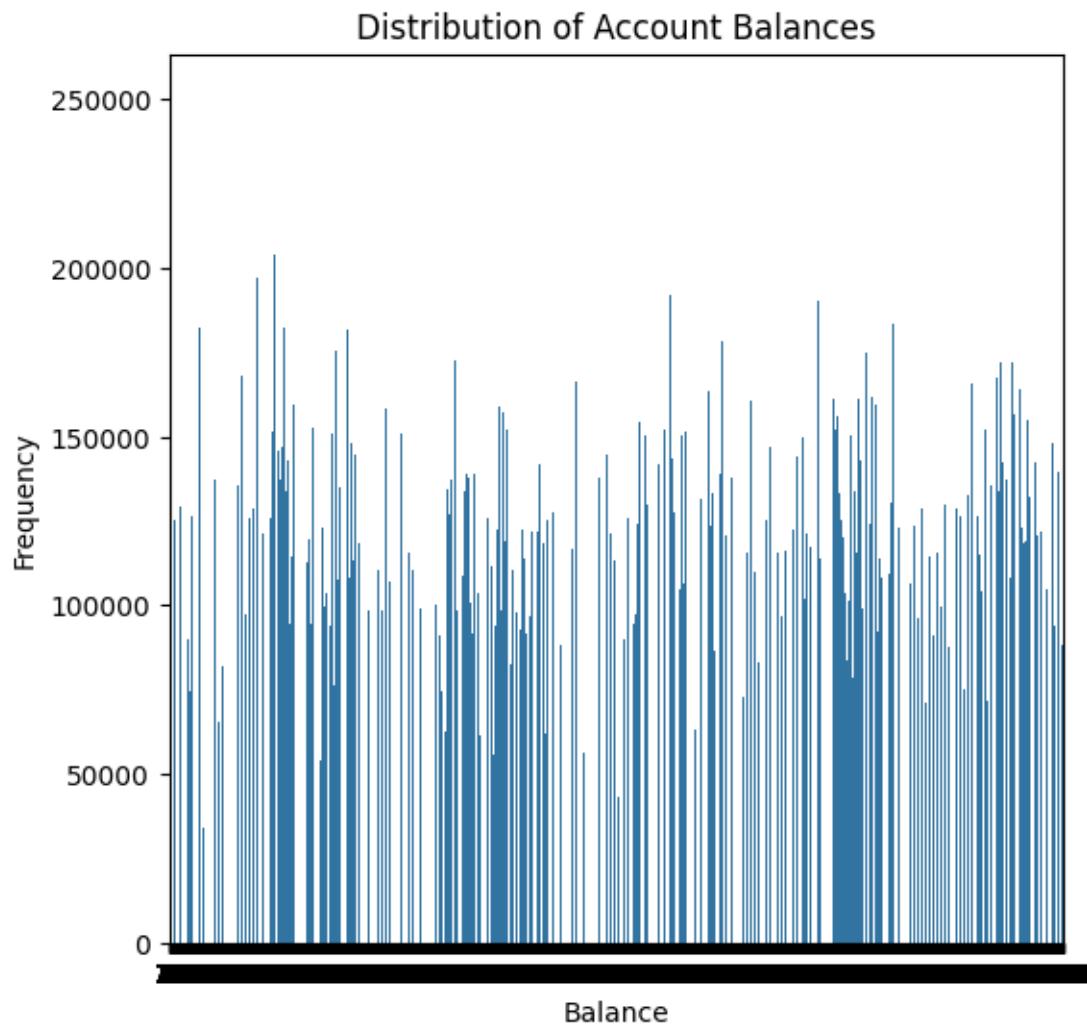


```
[15]: plt.figure(figsize=(6,6))

sns.barplot(Dataset["Balance"])

plt.xlabel("Balance")
plt.ylabel("Frequency")
plt.title("Distribution of Account Balances")

plt.show()
```

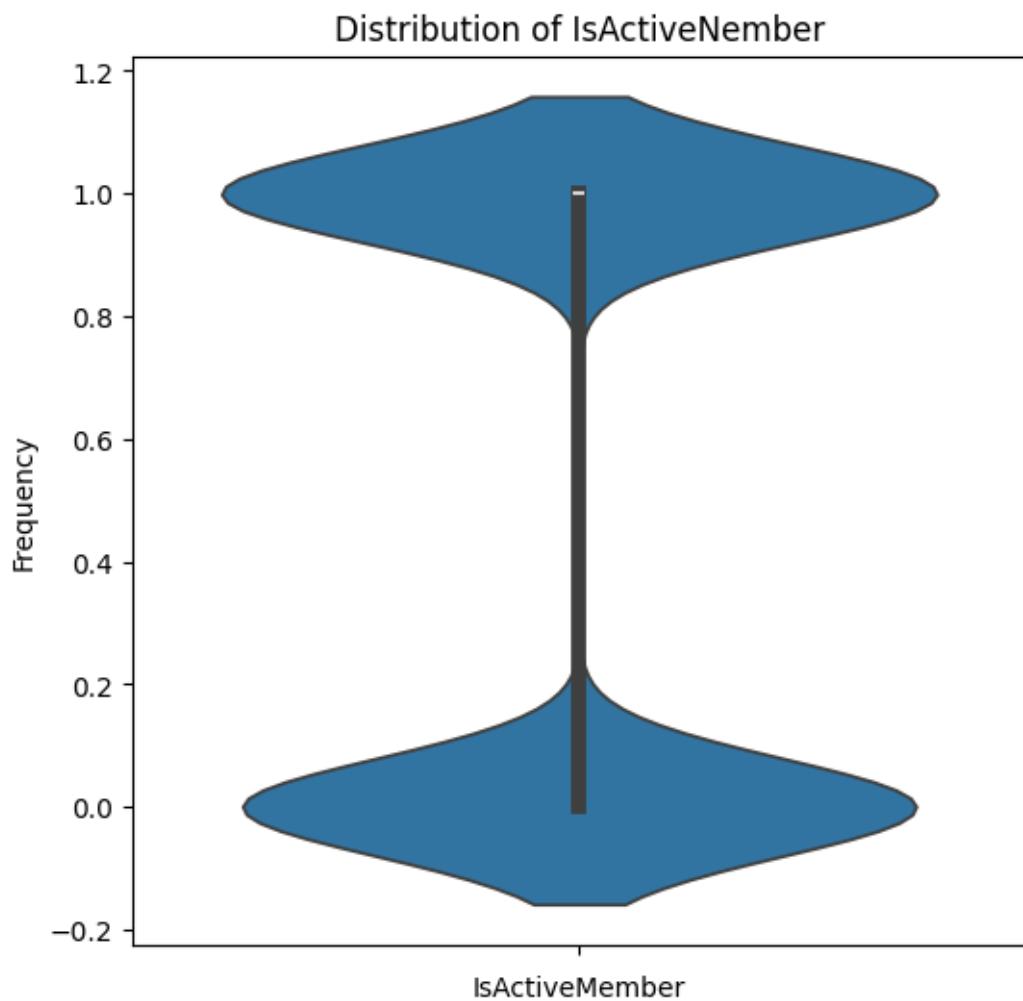


```
[16]: plt.figure(figsize=(6,6))

sns.violinplot(Dataset["IsActiveMember"])

plt.xlabel("IsActiveMember")
plt.ylabel("Frequency")
plt.title("Distribution of IsActiveMember")

plt.show()
```

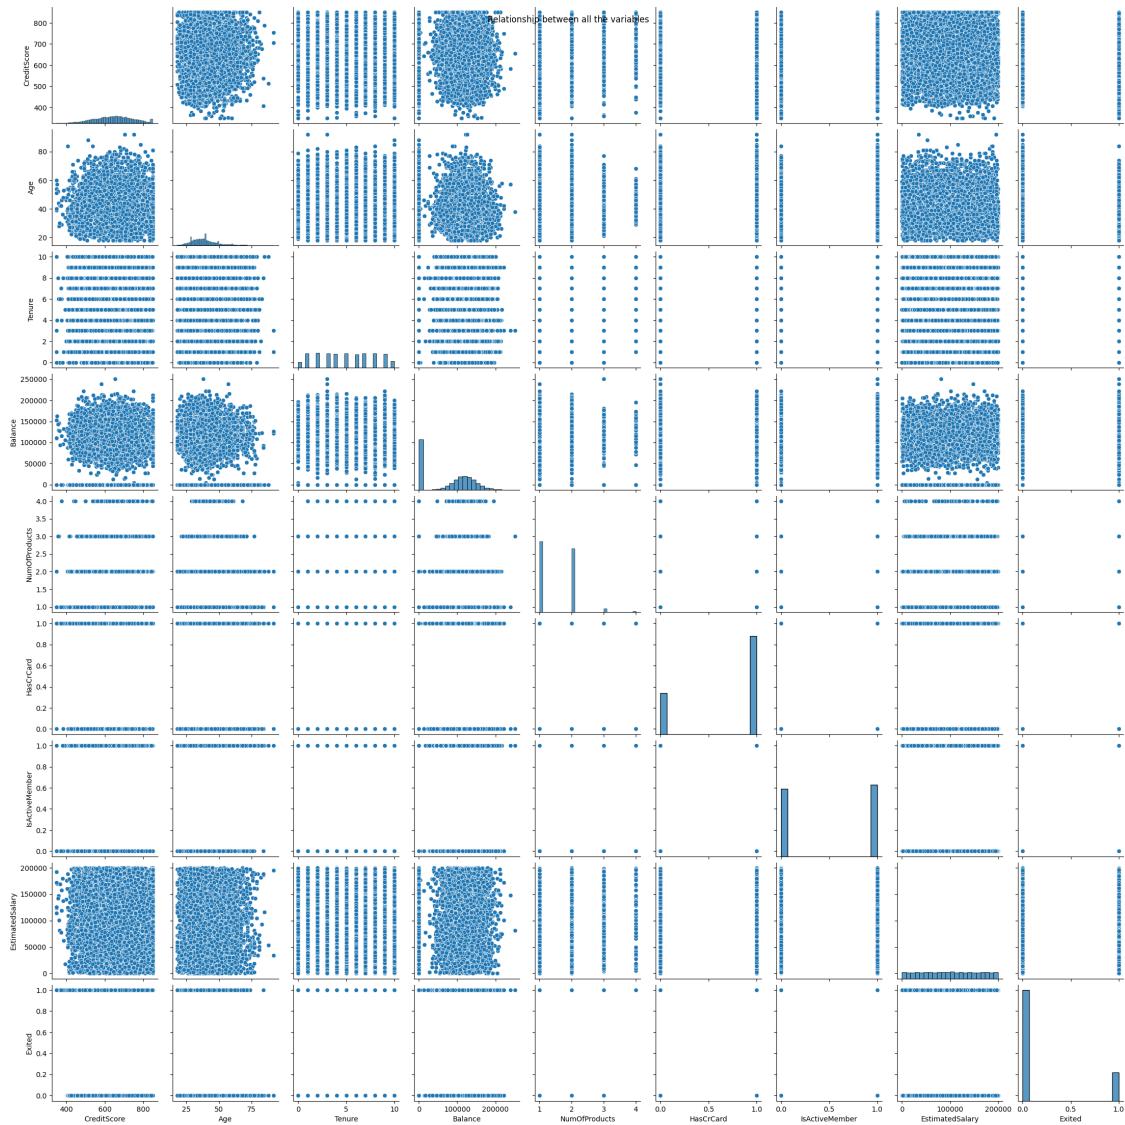


```
[17]: plt.figure(figsize=(6,6))

sns.pairplot(Dataset)

plt.suptitle("Relationship between all the variables")
```

```
[17]: Text(0.5, 0.98, 'Relationship between all the variables')
```



Clean and prepare the Dataset.

```
[18]: # Import Libraries.

import pandas as pd
import numpy as np

import matplotlib as pyplot
from matplotlib import pyplot as plt
%matplotlib inline

import seaborn as sns
```

```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import LabelEncoder

from sklearn.metrics import accuracy_score, confusion_matrix

```

[19]: `print(Dataset.isnull().sum())`

```

CreditScore      0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts   0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64

```

[20]: `dataset = Dataset.dropna()`

[21]: `dataset = pd.get_dummies(Dataset, drop_first=True)`

[22]: `dataset.head()`

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
0	619	42	2	0.00		1	1
1	608	41	1	83807.86		1	0
2	502	42	8	159660.80		3	1
3	699	39	1	0.00		2	0
4	850	43	2	125510.82		1	1

	IsActiveMember	EstimatedSalary	Exited	Geography_Germany	\
0	1	101348.88	1	False	
1	1	112542.58	0	False	
2	0	113931.57	1	False	
3	0	93826.63	0	False	
4	1	79084.10	0	False	

	Geography_Spain	Gender_Male
0	False	False
1	True	False
2	False	False
3	False	False

```
4           True        False
```

Encode the categorical features such as geography and Gender.

One Hot Encoding.

```
[23]: dataset[["Geography_Germany", "Geography_Spain", "Gender_Male"]].head()
```

```
[23]:    Geography_Germany  Geography_Spain  Gender_Male
0            False          False        False
1            False          True         False
2            False          False        False
3            False          False        False
4            False          True         False
```

```
[24]: # One-hot Encoding For the Geography and Gender.
```

```
dataset = pd.get_dummies(dataset,
                         columns=["Geography_Germany", "Geography_Spain", "Gender_Male"],
                         drop_first=True)
```

```
[25]: dataset.head()
```

```
[25]:    CreditScore  Age  Tenure      Balance  NumOfProducts  HasCrCard  \
0       619     42      2      0.00             1            1
1       608     41      1   83807.86             1            0
2       502     42      8  159660.80             3            1
3       699     39      1      0.00             2            0
4       850     43      2  125510.82             1            1

   IsActiveMember  EstimatedSalary  Exited  Geography_Germany_True  \
0              1        101348.88      1                False
1              1        112542.58      0                False
2              0        113931.57      1                False
3              0        93826.63      0                False
4              1        79084.10      0                False

   Geography_Spain_True  Gender_Male_True
0            False            False
1            True            False
2            False            False
3            False            False
4            True            False
```

Label Encoding.

```
[26]: # Initialize Label Encoder.
le = LabelEncoder()
```

```

# Encode Geography_Spain_True.
dataset["Geography_Spain_True"] = le.
    ↪fit_transform(dataset["Geography_Spain_True"])

# Encode Geography_Germany.
dataset["Geography_Germany_True"] = le.
    ↪fit_transform(dataset["Geography_Germany_True"])

# Encode Gender_Male_True.
dataset["Gender_Male_True"] = le.fit_transform(dataset["Gender_Male_True"])

```

[27]: dataset.head()

```

[27]:   CreditScore  Age  Tenure      Balance  NumOfProducts  HasCrCard  \
0          619    42       2        0.00           1            1
1          608    41       1    83807.86           1            0
2          502    42       8   159660.80           3            1
3          699    39       1        0.00           2            0
4          850    43       2   125510.82           1            1

   IsActiveMember  EstimatedSalary  Exited  Geography_Germany_True  \
0                  1        101348.88     1                  0
1                  1        112542.58     0                  0
2                  0        113931.57     1                  0
3                  0        93826.63      0                  0
4                  1        79084.10      0                  0

   Geography_Spain_True  Gender_Male_True
0                      0                      0
1                      1                      0
2                      0                      0
3                      0                      0
4                      1                      0

```

Train a classification Model.

Logistic Regression.

[28]: x = dataset.drop("Exited", axis=1)  
y = dataset["Exited"]

[29]: x = pd.get\_dummies(x, drop\_first=True)

[30]: x\_train, x\_test, y\_train, y\_test = train\_test\_split(x, y, test\_size=0.2, ↪random\_state=42)

[31]: print("Training Data:", x\_train.shape)  
print("Testing Data:", x\_test.shape)

```
Training Data: (8000, 11)
Testing Data: (2000, 11)
```

```
[32]: logistic_model = LogisticRegression(max_iter=1000)
logistic_model.fit(x_train, y_train)

y_predict_logistic = logistic_model.predict(x_test)
```

```
c:\Users\zaigh\AppData\Local\Programs\Python\Python313\Lib\site-
packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Evaluate the Model using the Confusion Matrix and the Accuracy.

```
[33]: print("Logistic Regression Accuracy: \n",
          accuracy_score(y_test, y_predict_logistic))

print("Confusion Matrix: \n",
      confusion_matrix(y_test, y_predict_logistic))
```

```
Logistic Regression Accuracy:
```

```
0.8145
```

```
Confusion Matrix:
```

```
[[1547  60]
 [ 311  82]]
```

```
[34]: cm = confusion_matrix(y_test, y_predict_logistic)
TN, FP, FN, TP = cm.ravel()

# Create a Structured Table.
metrics_table = pd.DataFrame({
    'Metric': ['TP', 'TN', 'FP', 'FN', 'Accuracy', 'Precision', 'Recall'],
    'Value': [TP, TN, FP, FN,
              (TP+TN)/(TP+TN+FP+FN),
              TP/(TP+FP),
              TP/(TP+FN)]})

print(metrics_table)
```

	Metric	Value
0	TP	82.000000

```

1      TN  1547.000000
2      FP   60.000000
3      FN  311.000000
4  Accuracy   0.814500
5  Precision   0.577465
6  Recall     0.208651

```

Decision Tree.

```
[35]: decisionTree_model = DecisionTreeClassifier(random_state=42)

decisionTree_model.fit(x_train, y_train)

y_predict_decisionTree = decisionTree_model.predict(x_test)
```

Evaluate the Model using the Confusion Matrix and the Accuracy.

```
[36]: print("Decision Tree Accuracy: \n",
           accuracy_score(y_test, y_predict_decisionTree))

print("Confusion Matrix: \n",
      confusion_matrix(y_test, y_predict_decisionTree))
```

Decision Tree Accuracy:

0.7805

Confusion Matrix:

```
[[1362 245]
 [ 194 199]]
```

```
[37]: cm = confusion_matrix(y_test, y_predict_decisionTree)
TN, FP, FN, TP = cm.ravel()

# Create a Structured Table.
metrics_table = pd.DataFrame({
    'Metric': ['TP', 'TN', 'FP', 'FN', 'Accuracy', 'Precision', 'Recall'],
    'Value': [TP, TN, FP, FN,
              (TP+TN)/(TP+TN+FP+FN),
              TP/(TP+FP),
              TP/(TP+FN)]})
print(metrics_table)
```

	Metric	Value
0	TP	199.000000
1	TN	1362.000000
2	FP	245.000000
3	FN	194.000000
4	Accuracy	0.780500
5	Precision	0.448198
6	Recall	0.506361

Analyze feature importance to understand what influences churn.

Skills: Categorical data encoding (Label Encoding/One-Hot Encoding). Supervised classification modeling. Understanding and interpreting feature importance.

Task completed. Best Wishes. Zaigham Abbas.